



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

Processor Design
(PD)

Module Implementation

3rd Report

Jordi Solà, Ying hao Xu

November 14, 2018

Contents

1	Specification update	2
1.1	Feed array	2
1.2	Drain array	3
1.3	Systolic Array Wrapper	4
2	Development status	4
2.1	Infrastructure	4
2.1.1	Test tools status	5
2.1.1.1	Structure of the system	5
2.1.1.2	Installation and environment setup	5
2.1.1.3	Specification of a testbench and compilation	5
2.1.1.4	Execution	6
2.1.1.5	Viewing the trace (<i>unstable</i>)	6
2.1.2	Next steps	7
2.2	Model	7
2.2.1	HDL modules	7
2.2.2	Next steps	8
3	Validation	8
3.1	Testplan	8
3.1.1	PE	8
3.1.2	Systolic array	8
3.1.3	Wrapper	9
3.2	Next steps	9
4	Verification	9
4.1	Test specification	9
4.1.1	Hardware	9
4.1.2	Input	9
4.1.3	Reproducibility	10
4.2	Analysis	10
4.2.1	Processing Element	10
4.2.2	Drain channel	10
4.2.3	Wrapper	10
	References	12
A	Appendices	12
A.1	Dockerfile	12
A.2	Environment script: <code>env.sh</code>	13
A.3	Compile script: <code>compile</code>	13
A.4	Generic main file: <code>main.cpp</code>	14
A.5	Run script: <code>run</code>	16
A.6	View script: <code>view</code>	17

1 Specification update

In the previous report, we shown the structure of a single PE (Processing Element) and how they are interconnected in order to form the actual core of the systolic array. Something we did not cover in the last report is how this systolic array is going to be feed or drained. For this task, we assumed that our internal cache memory of the accelerator has a limited number of read / write ports.

1.1 Feed array

As the Systolic Array has to be feed following a specific pace, we decided to add what we call a feed array. This array is basically an array of flip-flops connected such in a way that each i element is delayed exactly one cycle after the previous element ($i - 1$). This mechanism lets us feed N elements into the Systolic Array in the same cycle, without having to worry about the rate in which the Systolic Array expects the data to enter.

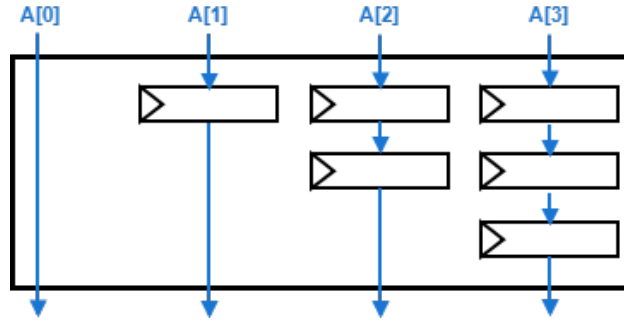


Figure 1: Feed array for $N = 4$

The basic structure can be seen in the Figure 1, in this case, as the HDL language we chose supports the **for** syntax, we defined our feed array as a macro that automatically generates the correct number of delay flip-flops given a **SIZE** parameter (Figure 2).

```
1  'define DELAY_ARRAY(CLK, RESET, EN, SIZE, ARRAY_DATA_I, ARRAY_DATA_O) \  
2      for (genvar gv_i=0; gv_i < SIZE; gv_i++) begin \  
3          if (gv_i == 0) begin\  
4              assign ARRAY_DATA_O[0] = (RESET == 1'b1) ? '0 : ARRAY_DATA_I[0];\  
5          end\  
6          else begin\  
7              logic [$bits(ARRAY_DATA_I[0])-1:0] delayer [gv_i:0];\  
8              for (genvar gv_j=0; gv_j < gv_i; gv_j++)\  
9                  'FF_RESET_EN(CLK, RESET, EN, delayer[gv_j], delayer[gv_j+1], '0)\  
10             assign delayer[0] = ARRAY_DATA_I[gv_i];\  
11             assign ARRAY_DATA_O[gv_i] = delayer[gv_i];\  
12         end\  
13     end
```

Figure 2: Delay array macro in SystemVerilog

The number of required flip-flops given an n can be expressed using the following formula:

$$A(n) = \frac{n^2 - n}{2}$$

1.2 Drain array

Similarly as in the feed array, our Systolic Array design does not outputs all the data in the same cycle. Furthermore, as we are using a shared drain channel between two consecutive columns, it is critical to read and store the data in the exact cycle. For this task, we followed a similar approach as in the feed array by using a delay array in order to abstract to the other functional units, the pace in which the results are available.

In our systolic array, the order in which the results flow out through the drain channel can be seen in Table 1. Each element of the table represents in which cycle it is going to be available on the output ports. For example, for the element (0,0), we are considering that it is going to be available in the cycle 1, the element (0,1) in the cycle 2 and so on. Notice that the elements flow out in a interleaved fashion, so for example, in the cycle 3 and 4, the drain channel outputs elements from the first and the second row.

1		2	3		4
3		4	5		6
5		6	7		8
7		8	9		10

Table 1: Systolic Array drain order for $N = 4$. Each cell represents the cycle when it is available.

As the drain channel is half the size of the systolic array, consecutive elements are available within one cycle of distance and they will flow out from the same output port. When we were thinking about this unit, we came up with two different approaches (Figure 3a and Figure 3b).

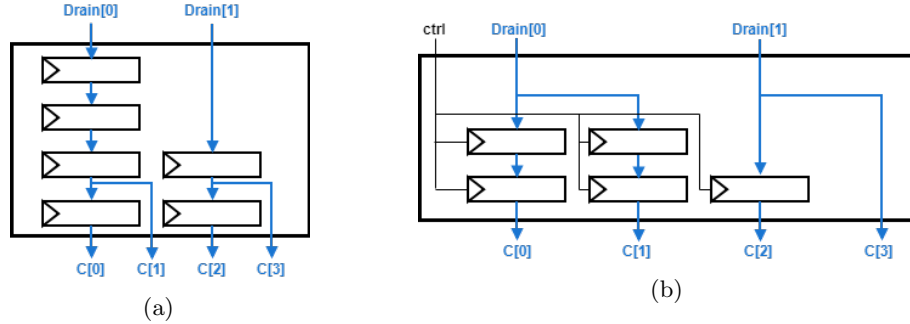


Figure 3: Drain array proposals for $N = 4$

Figure 3a shows the most natural way to design the delay array. In this case, each element is processed sequentially and during the last stage, the element in the previous stage is bypassed, so we can read the N result elements at once.

Figure 3b focuses on reducing the number of flip-flop toggles. In this design, the delay array granularity is a pair (two elements) and only when both elements are ready, the data is passed to the next stage. This behavior forced us to use a control signal (**ctrl**) that must be in sync with the rest of the system in order to write the data to the correct flip-flop.

If we compare Figure 3a and Figure 3b, the number of flip-flops needed in both cases is almost the same. However, for the second one, as we are passing the elements to the next stage every two cycles, the number of toggles is significantly reduced. This can be easily seen if we analyze it from the point of view of a single element. While in the Figure 3a approach, in the worst case an element has to traverse N flip-flops, in the Figure 3b approach, also in the worst case, the number of flip-flops to traverse is reduced by half ($N/2$), as the elements are passed to the next stage every two cycles instead of every cycle.

As power is something important in our design, even if the final design becomes more complex, we decided to go with the Figure 3b approach, as it requires almost the same number of flip-flops as the Figure 3a approach, but consuming half of the dynamic power. The number of flip-flops that are needed for the 3b approach can be expressed using the following formula:

$$D(n) = \frac{n^2 + 2n}{4} - 1$$

1.3 Systolic Array Wrapper

Now that we have defined the feed array and the drain array, we joined them and created the Systolic Array Wrapper module. This module is basically the Systolic Array plus 2 feed arrays (as we have two inputs) and a single drain array (as we only have one output). Figure 4 shows how the systolic array is going to be connected to the new arrays and which input/output ports will be exposed to the rest of components.

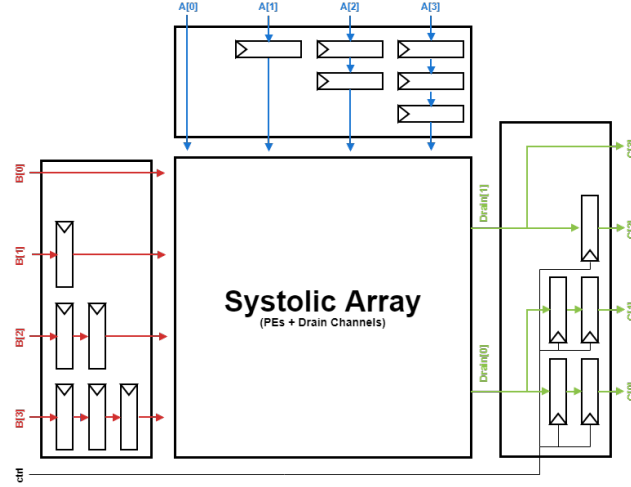


Figure 4: Systolic Array wrap module

This new wrap module let us forget about the pace in which the systolic array read / writes the data, simplifying the interaction between the systolic array and the other units. Considering that the wrap module has 2 feed arrays and 1 drain array, the total number of extra flip-flops that we need is given by the addition of the aforementioned elements, thus it grows lineally with the array area:

$$T(n) = 2 \cdot A(n) + D(n) = 2 \cdot \frac{n^2 - n}{2} + \frac{n^2 + 2n}{4} - 1 = \frac{5}{4} \cdot n^2 - \frac{1}{2} \cdot n - 1$$

2 Development status

2.1 Infrastructure

One of our main goals for the project development was to build and maintain an infrastructure that allow us to automatize the tests and perform some kind of validation of the project.

This implies two main tasks, first, we need to develop some testing tools allowing us to compile and execute our testbenches, and second, we need to design a testplan and build the corresponding tests for each component.

2.1.1 Test tools status

2.1.1.1 Structure of the system

We wanted the system to be easily reproducible without a complicated setup in as many diverse environments as possible, and this conditioned our design decisions.

Given that, we decided to use Linux containers (in particular Docker) in order to isolate our infrastructure environment, thus allowing us to run out-of-the-box, independently of the host environment. There is a lot of literature on the benefits of containerization on reproducible research, and one of the easiest to understand is Carl Boettiger’s introduction [1].

Keeping that on mind we created a Docker image containing all our software dependencies. All our scripts spawn an instance of the container mounting our project folder and doing its job inside our isolated environment.

This image is published at the docker store under `jsola/verilator` and can be imported with the following command:

```
1 docker pull jsola/verilator
```

The Dockerfile can be found in the Appendix A.1.

2.1.1.2 Installation and environment setup

We wanted to automatize as much as possible the setup, thus we created a script to automatically detect and configure the environment variables needed to run our tests and pulling our docker image as well. This script can be found at the Appendix A.2.

Thus, given that we have Docker installed on our machine, the only thing needed to setup the environment for the project is to clone it from github and source this script:

```
1 git clone https://github.com/ying27/matrix-multiplication-accelerator.git
2 cd matrix-multiplication-accelerator
3 source env.sh
```

2.1.1.3 Specification of a testbench and compilation

In order to perform tests of specific parts of the design, we need to be able to create different testbenches instantiating the top module of the hierarchy we want to test.

In order to simplify the compilation flow, we also created a script that given a parameter specifying the name of the testbench, it automatically looks for the testbench top at the path `tb/TB_NAME/TB_NAME.sv` and compiles the model.

One extra parameter to consider when compiling the model is whether we want to see a trace of the result or not. One could think that it makes no harm to always compile it with trace support, but it adds some specific constraints to the compiler resulting both in a larger compilation and execution time. Therefore our script has an optional parameter `-trace` used to specify that we want a model with support for traces.

Also, some specific testbenches might use extra files to execute (due to `$fopen` or `$freadmem` in the code), thus, our tool lets us specify a test name using `-test` and all the files under the directory are linked to the future build folder. This parameter looks for all the files inside the folder `tb/TESTBENCH/test/TESTNAME`.

With this, in order to create a testbench of a specific element, one can simply create the file `element_tb.sv` in the folder `tb/element_tb` with the following baseline testbench:

```

1  `ifndef TIMEOUT
2      `define TIMEOUT 100
3  `endif
4
5  module element_tb(
6      input clk,
7      input reset
8  );
9
10     element element_name(
11         .clk_i(clk),
12         .rst_i(reset),
13         .* //Automatically capture the defined signals
14     );
15
16     //Timeout control
17     int cycle_counter;
18     `FF_RESET(clk, reset, cycle_counter+1, cycle_counter, '0);
19
20     always_comb begin
21         if(cycle_counter == `TIMEOUT) $finish;
22     end
23
24 endmodule

```

To compile the testbench, we have to execute the following command:

```
1 compile -top element_tb
```

Or, if we want to be able to produce a trace afterwards:

```
1 compile -top element_tb -trace
```

This will perform the aforementioned flow and create a folder with the build at **tb/element_tb/build**.

Both, the compile script as well as the generic **main.cpp** file for verilator can be found at appendices A.3 and A.4.

2.1.1.4 Execution

Following the criteria we've seen before, we also have created a script in order to simplify the execution of the model. This script has one mandatory argument which is the name of the testbench that needs to be run. If the testbench was compiled for tracing we can pass as an argument **--trace target.flt** to produce a tracefile. The script can be found at appendix A.5.

Once the test execution finishes, a soft-link at **\$REPOROOT/last_run** pointing to the execution directory will be created. So, in order to get a trace of our previous model **element_tb** we do it like this:

```
1 run element_tb --trace trace.flt
```

2.1.1.5 Viewing the trace (*unstable*)

There's a tool in development to be able to launch gtkwave inside the docker environment in order to make it generic too, however, as forwarding the X display to the host is not host independent, and so far this only works on OSX with XQuartz 2.7 and it's still unstable. The script can be found at the appendix A.6.

This can be done by using:

```
1 view trace.flt
```

2.1.2 Next steps

Given our current status, the following steps would be to further automatize the testing and consolidate a standard for test specification. This implies also to build the necessary structures for applying the correspondent stimulus to the simulations, making them UVM compliant if possible.

2.2 Model

One of the main goals of our project is to develop HDL that is scalable and easy to understand. For this reason, we have defined some key parameters that can be found in `inc/common_pkg.sv`:

- **DATA_WIDTH:** This defines the width of the data that the systolic array operates with.
- **SYS_ARRAY_SIZE:** This defines the number of rows/columns of the systolic array.
- **ADDR_WIDTH:** This is the width of the memory addresses used to read/write data into the accelerator.

Following the same idea, all the data types are well defined inside the `inc/common_pkg.sv` as a `struct packed` type, which let us make the design even more generic and easy to modify.

As one of the potential extensions is to apply clock gating to the flip-flops, we have defined 3 macros that mimics the behaviour of a flip-flop:

- `FF_RESET(CLK, RESET, DATA_I, DATA_O, DEFAULT)`
- `FF_EN(CLK, RESET, EN, DATA_I, DATA_O, DEFAULT)`
- `FF_RESET_EN(CLK, RESET, EN, DATA_I, DATA_O, DEFAULT)`

By using these macros, we can easily change the flip-flops to a gated one but it also makes the files easier to read and more close to the actual implementation. These macros can be found in the file `inc/common.svh`

2.2.1 HDL modules

Currently, the core of the systolic array is working. Thus is a module that is capable of successfully complete a matrix matrix multiplication. In order to accomplish this task, we have implemented the following modules:

- **PE:** Processing Element that computes the MAC (Multiply Accumulate) operation (`pe.sv`).
- **Drain channel:** implementation of a single drain channel that connects two PEs from the same row (`drain_channel.sv`).
- **Systolic array:** module that instantiates the PEs and the drain channels given a size and interconnects them properly (`systolic_array.sv`).
- **Drain array:** implementation of the drain array described in the previous section (`drain_array.sv`).
- **Systolic array wrap:** module that connects the systolic array and the feed/drain arrays (`systolic_array_wrap.sv`).

However, the control logic is not ready yet and the signals must be carefully driven by a behavioural model.

2.2.2 Next steps

As mentioned above, we currently have implemented the core of the accelerator. As part of the next steps, we will have to implement the controller. This unit will be responsible of accessing to the accelerator's cache and feed the systolic with the desired data. On the other hand, as the controller knows at any time the state of the systolic array, it will also be responsible of driving the **last** and **ctrl** signals.

Once we have a controller, a cache will be implemented too. This will store both, the data required for feeding the accelerator and the results of the matrix matrix multiplication. As the project main focus is not the memory subsystem, for this task we will assume an ideal memory model.

And last but not least, an instruction decoder will be required too. This unit will interpret a mini-ISA (Instruction Set Architecture) that follows the RISC (Reduced Instruction Set Computer) style. The instructions it will cover are described in the previous report, but basically are read, write and operate instructions.

3 Validation

"Validation. The assurance that a product, service, or system meets the needs of the customer and other identified stakeholders. It often involves acceptance and suitability with external customers. Contrast with verification."[2]

3.1 Testplan

In this section we will try to describe a testplan to verify that our DUTs (Design Under Test) works properly, following the expected behaviour.

3.1.1 PE

For the PE, we intend to test the following situations:

	Test	Status
1	Send a request to a PE without the last bit set.	PASSING
2	Send a request to a PE with the last bit set.	PASSING
3	Send a sequence or requests to a PE without any request with the last bits set.	PASSING
4	Send a sequence of requests to a PE finishing it with a request with the last bit set.	PASSING
5	Send it a sequence of requests with interleaved requests with the last bit set.	PASSING
6	Send it a request which produces a multiplication overflow.	PENDING
7	Send it a sequence of request which produce an addition overflow.	PENDING
8	Send it a request with the last bit set only on one of the inputs.	PENDING

3.1.2 Systolic array

This are the ones we though could be good for our array model:

	Test	Status
1	Send a request through each of the channels without the last bit set.	PASSING
2	Send a request through each of the channels with the last bit set.	PASSING
3	Send a sequence of requests through each channel without any of them having the last bit set.	PASSING
4	Send a sequence of requests through each channel with the last of them having the last bit set.	PASSING
5	Send a sequence of requests through each channel with interleaved requests with the last bit set.	PASSING

3.1.3 Wrapper

This are the tests considered to be run for the wrapper module:

	Test	Status
1	Send a request to the wrapper without the last bit set.	PASSING
2	Send a request to the wrapper with the last bit set.	PASSING
3	Send a sequence of requests to the wrapper without any of them having the last bit set.	PASSING
4	Send a sequence of requests to the wrapper with the last of them having the last bit set.	PASSING
5	Send a sequence of requests to the wrapper with interleaved requests with the last bit set.	PASSING

3.2 Next steps

On the validation side, this are the next steps planned for the next release:

- Creation of checkers for each relevant module following UVM standards.
- Extension of the testing environment taking into account the feedback provided by the checker, detecting execution errors.
- Addition of assertions for explicit implementation error detection.
- Designing the testplan for the forthcoming modules, and do an intense revision and possible extension of the current ones.
- Implement and monitor line coverage, toggle coverage, and functional coverage metrics.

4 Verification

"Verification. The evaluation of whether or not a product, service, or system complies with a regulation, requirement, specification, or imposed condition. It is often an internal process. Contrast with validation."[2]

4.1 Test specification

4.1.1 Hardware

In order to simplify the analysis of the trace while keeping some complexity, we decided to test the model with the following configuration:

- Data width of 1 byte.
- Array of 4x4 PEs.

As we are not using the memory yet, the value for the address length is irrelevant.

4.1.2 Input

In order to verify the behaviour of our build model, we have performed a simple matrix multiplication:

$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 8 & 7 & 6 \\ 5 & 4 & 3 & 2 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 8 & 7 & 6 & 5 \\ 4 & 3 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 51 & 47 & 43 & 39 \\ 123 & 119 & 115 & 111 \\ 129 & 133 & 137 & 141 \\ 57 & 61 & 65 & 69 \end{bmatrix}$$

In order to do so, we had to add to the testbench the capacity to read values from a pre-formatted file. This was achieved by adding several `$scanf`'s to the code and creating a file with the data we wanted to feed at `tb/array_wrap_tb/test/base`. The test we created feeds the array two times with the aforementioned matrices, in order to see the interaction of several operations.

4.1.3 Reproducibility

Our test can be reproduced by doing the following:

```
1 git clone https://github.com/ying27/matrix-multiplication-accelerator.git
2 cd matrix-multiplication-accelerator
3 git checkout 85f4337f9e83f69d2ce303fee20c6891757cef91
4 source env.sh
5 compile -top array_wrap_tb -trace -test base
6 run array_wrap_tb --timeout 10000 --trace test.fst
```

This will produce the trace file `test.fst` at `$REPOROOT/last_run/test.fst` which can be later opened with `gtkwave`. The signal files used to see the screen captures can be found at `$REPOROOT/tb/array_wrap_tb/profiles`.

4.2 Analysis

4.2.1 Processing Element

First, let us take a look at the behaviour of the most basic unit, the PE. For this, we've picked as an example the first element of the array. Each PE is in charge of computing the element of the resultant matrix with its same indices, thus, each PE is in charge of computing the following operation:

$$value_{i,j} = \sum_{k=0}^N a_{i,k} * b_{k,j}$$

The PE structure is targeted to perform an accumulate each of the products each consecutive cycle. With this, after running the test, Figure 6 shows the signals of $PE_{0,0}$. There, we can see that the PE is correctly feed, receiving each value of the matrix at the specified cycle and performing a multiply plus addition operation, accumulating the result in `result_n` (which is the name for the wire feeding the partial register). We can see how the element starts to get feed at marker A which causes it to start performing computation and storing partial results, and receives its `last` control signal at marker B, which produces the emission of the results to the drain channel and the resetting of the partial stored value.

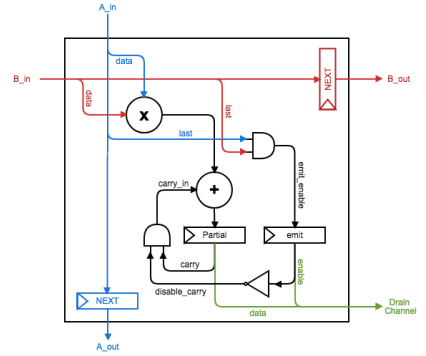


Figure 5: Reminder of the internal structure of a PE

4.2.2 Drain channel

As for the Drain Channel, looking at the trace for the first drain channel (Figure 7) we can see how the PE's finish their computations in a way that they do not collide with each other and leave the pipeline as specified in Section 1.2.

4.2.3 Wrapper

For the top level, our feed and drain arrays allow us to feed each line of data at once, as well as retrieve the results also line by line. On Figure 8 we can see a trace of the behaviour of our wrapper. First of all, we can see how it is feed, line by line in case of input A, and column by column in case of B, and how we retrieve it line by line at output C. We can see the first operation being fed between markers A and B, and how the result of that operation is emitted every two cycles after marker C until marker D.




Figure 6: PE trace during execution




Figure 7: Drain channel trace during execution

Also, we can observe the delay we have for issuing the second operation. This will delay the emission of the result in order to avoid interfering with the previous operation's result, right as we predicted in our last assignment.

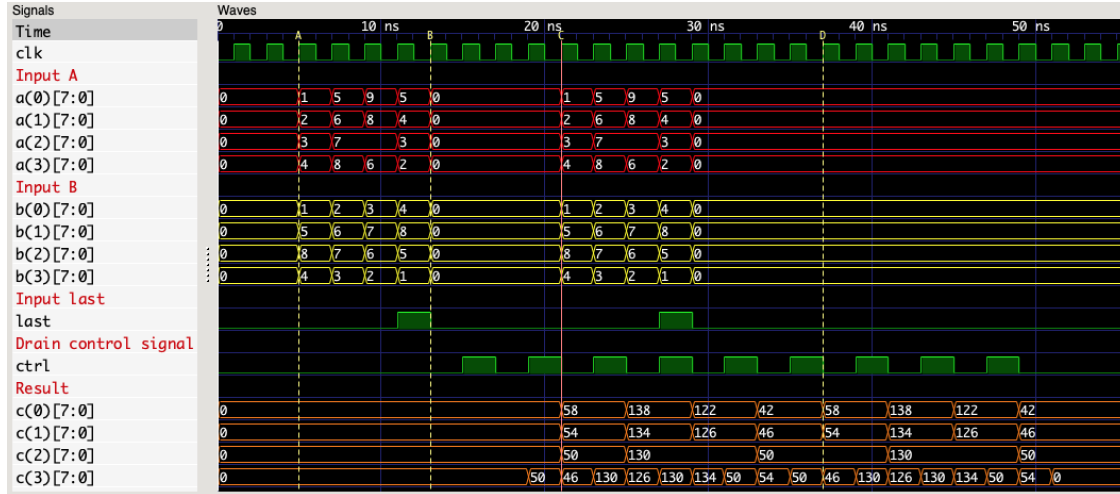


Figure 8: Systolic Array wrap module trace during execution.

References

- [1] Carl Boettiger. *An introduction to Docker for reproducible research*. Tech. rep.
- [2] *IEEE P1490/D1, May 2011 - IEEE Draft Guide: Adoption of the Project Management Institute (PMI) Standard: A Guide to the Project Management Body of Knowledge (PMBOK Guide)-2008 (4th edition)*. IEEE, 2011. ISBN: 9780738167886. URL: <https://ieeexplore-ieee-org.recurtos.biblioteca.upc.edu/document/5937011>.

A Appendices

A.1 Dockerfile

```

1 FROM debian:stretch
2
3 WORKDIR /tmp
4 RUN apt-get -y update
5 RUN apt-get -y install apt-utils
6
7 #Installing pre-requisites
8
9 RUN apt-get -y install zlib1g-dev bash build-essential git make
10 RUN apt-get -y install autoconf perl gcc g++ flex bison
11
12 # Installation of Verilator
13
14 RUN git clone http://git.veripool.org/git/verilator
15 WORKDIR /tmp/verilator
16
17 RUN git pull
18 RUN git checkout stable          # Use most recent release
19 RUN autoconf
20 RUN ./configure
21 RUN make -j $(nproc)

```

```

22 RUN make install -j $(nproc)
23 WORKDIR /tmp
24
25 # Clean installation files
26 RUN rm -rf verilator
27
28 # Installation of gtkwave
29 RUN apt-get -y install gtkwave
30
31 # Adding the installed software to the path of the container
32 ENV PATH /usr/local/bin:$PATH
33 ENV SHELL /bin/bash

```

A.2 Environment script: **env.sh**

```

1 #Detect the path to the top of the repo from which we are calling the script
2 export REPOROOT=$( cd $(dirname "$0"); pwd -P)
3
4 #Add our scripts to the path
5 if [ -z ${OLD_PATH+x} ]; then
6     export OLD_PATH=$PATH;
7 else
8     export PATH=$OLD_PATH;
9 fi
10 export PATH="$REPOROOT/tools/scripts:$PATH"
11
12 #Specify the path of the rtl and testbenches
13 export RTLROOT="$REPOROOT/src"
14 export TBROOT="$REPOROOT/tb"
15
16 #Pull docker image if necessary
17 export DOCKERTAG="jsola/verilator:latest"
18 if [[ "$(docker images -q $DOCKERTAG 2>/dev/null)" == "" ]]; then
19     docker pull $DOCKERTAG
20 fi

```

A.3 Compile script: **compile**

```

1 #!/bin/bash
2
3 #Variable specification
4 export RTL_TOP
5 export DOCKER_NAME=compile
6 export EXTRA_VERLIATOR_ARGS
7 export CFLAGS
8
9 #Detect arguments
10 while [ $# -ne 0 ]
11 do
12     arg="$1"
13     case "$arg" in
14         -top) #TB path
15             RTL_TOP=$2
16             CFLAGS+=\ -DCXX_TOP_LEVEL=V$2
17             shift

```

```

18         ;;
19         -trace)#Trace or not
20         EXTRA_VERILATOR_ARGS+=--trace-fst
21         CFLAGS+=\ -DTRACE
22         ;;
23     esac
24     shift
25 done
26
27 export TARGET_PATH=$TBROOT/$RTL_TOP/build
28 echo "RTL_TOP:_$RTL_TOP"
29
30 #Clean
31 echo "Cleaning_last_build"
32 rm -rf $TARGET_PATH
33
34 #Docker container creation
35 echo "Creating_compile_container"
36 docker run -w $REPOROOT -i -t -d -v $REPOROOT:$REPOROOT \
37 --name=$DOCKER_NAME $DOCKERTAG > /dev/null
38
39 #Verilog compilation
40 echo "Compiling_verilog"
41 compile_cmd="verilator -Wall $EXTRA_VERILATOR_ARGS --Mdir $TARGET_PATH -CFLAGS \
42 \"$CFLAGS\" --cc $RTL_TOP -f $REPOROOT/tools/conf/incdir.f -y $TBROOT/$RTL_TOP \
43 --exe $REPOROOT/tools/csrc/*.cpp"
44 echo $compile_cmd
45 docker exec -t -w $REPOROOT $DOCKER_NAME bash -c "$compile_cmd"
46
47 #C Compilation
48 echo "Compiling_C"
49 compile_cmd="OPT=-DVL_DEBUG make -j -C $TARGET_PATH -f V$RTL_TOP.mk V$RTL_TOP"
50 echo $compile_cmd
51 docker exec -t -w $REPOROOT $DOCKER_NAME bash -c "$compile_cmd"
52
53 #Cleaning the environment
54 echo "Removing_container"
55 docker stop $DOCKER_NAME > /dev/null
56 docker rm $DOCKER_NAME > /dev/null

```

A.4 Generic main file: main.cpp

```

1  #include "verilated.h"
2  #include <iostream>
3  using namespace std;
4  #include "verilated_fst_c.h"
5  #include <getopt.h>
6
7  #define PATH_LENGTH 1000
8
9  #define mStr(x) #x
10 #define mStr_(x) mStr(x)
11 #define topModuleInc mStr_(CXX_TOP_LEVEL.h)
12

```

```

13 #include topModuleInc
14
15 void usage(char* name) {
16     cerr << "Usage:_" << name << "_[--trace_traceFile.fst]_[--timeout_timeout]_[--help]"
17     cerr << "\tTrace:_off_by_default" << endl << "\tTimeout:_default_1" << endl;
18     exit(EXIT_FAILURE);
19 }
20
21 void getArgs(int argc, char* argv[], uint64_t &timeout, bool &trace, char* traceFile){
22
23     trace = false;
24     timeout = 1;
25
26     struct option argOptions[] = {
27         { "trace", required_argument, 0, 1},
28         { "timeout", required_argument, 0, 2},
29         { "help", no_argument, 0, 3}
30     };
31
32     int idx;
33     while ( (idx = getopt_long(argc, argv, "", argOptions, NULL)) != -1 ){
34         switch(idx){
35             case 1:
36                 trace = true;
37                 strncpy(traceFile, optarg, PATH_LENGTH);
38                 break;
39             case 2:
40                 timeout = strtol(optarg, NULL, 0);
41                 break;
42             default:
43                 usage(argv[0]);
44         }
45     }
46 }
47
48 int main(int argc, char** argv, char** env) {
49
50     Verilated::commandArgs(argc, argv);
51
52     bool trace;
53     char traceFile[PATH_LENGTH];
54     uint64_t timeout;
55
56     getArgs(argc, argv, timeout, trace, traceFile);
57
58     CXX_TOP_LEVEL* top = new CXX_TOP_LEVEL ();
59
60 #ifdef TRACE
61     VerilatedFstC* tfp = new VerilatedFstC;
62
63     if(trace){
64         Verilated::traceEverOn(true);
65         top->trace(tfp, 99);
66         tfp->open(traceFile);

```



```

67     }
68 #endif
69
70     int main_time = 0;
71     top->clk = 0;
72     top->reset = 1;
73
74 #ifdef TRACE
75     if(trace) tfp->dump(main_time);
76 #endif
77
78     top->eval();
79     while (!Verilated::gotFinish() && main_time < timeout) {
80         main_time += 1;
81
82         top->clk = 1;
83         top->eval();
84
85 #ifdef TRACE
86         if(trace) tfp->dump (main_time);
87 #endif
88
89         main_time += 1;
90         top->reset = 0;
91         top->clk = 0;
92         top->eval();
93
94 #ifdef TRACE
95         if(trace) tfp->dump (main_time);
96 #endif
97     }
98
99     if(main_time >= timeout) cout << "ERROR: timeout" << endl;
100
101 #ifdef TRACE
102     if(trace) tfp->close();
103 #endif
104     delete top;
105
106     exit(0);
107 }

```

A.5 Run script: **run**

```

1  #!/bin/bash
2
3  export DOCKER_NAME=run
4
5  export RTL_TOP=$1
6  shift
7  export TARGET_PATH=$TBROOT/$RTL_TOP/build
8
9  echo "Creating run container"
10 docker run -w $REPOROOT -i -t -d -v $REPOROOT:$REPOROOT --name=$DOCKER_NAME $DOCKERTAG :

```

```

11 echo "Running_Verilog"
12 run_cmd="$TARGET_PATH/V$RTL_TOP_.$*"
13 echo $run_cmd
14 docker exec -t -w $REPOROOT $DOCKER_NAME bash -c "$run_cmd"
15 echo "Removing_container"
16 docker stop $DOCKER_NAME > /dev/null
17 docker rm $DOCKER_NAME > /dev/null

```

A.6 View script: **view**

```

1  #!/bin/bash
2
3  export DOCKER_NAME=run
4
5  export RTL_TOP=$1
6  shift
7  export TARGET_PATH=$TBROOT/$RTL_TOP/build
8
9  echo "Creating_run_container"
10 docker run -w $REPOROOT -i -t -d -v $REPOROOT:$REPOROOT --name=$DOCKER_NAME $DOCKERTAG :
11 echo "Running_Verilog"
12 run_cmd="$TARGET_PATH/V$RTL_TOP_.$*"
13 echo $run_cmd
14 docker exec -t -w $REPOROOT $DOCKER_NAME bash -c "$run_cmd"
15 echo "Removing_container"
16 docker stop $DOCKER_NAME > /dev/null
17 docker rm $DOCKER_NAME > /dev/null

```